AI Prompt Log — Lunchline Partners Case Study

Aviv Sheriff

May 31, 2026

CONTENTS

Prompt Log — AI Usage Appendix

Running log of AI interactions for the Lunchline Partners case study.

Scope of this log

Reconstructed from the full Claude Code session transcripts for this project (**~524 user prompts across 2026-05-28 → 2026-06-01**, pulled via the Motif CLI). Sessions 1–6 are the detailed early entries (data pipeline, screening framework, memo system). **Sessions 7+ summarize the analytical core** — the multi-agent scoring calibration, its retirement in favour of the Deep Flow, the candidate tournament, the ARTW selection and build, and the deck iteration — which the earlier log did not capture.

Phase	Dates	What	Outcome
Infra & screening	5/28	Reproducible data pipeline; all-sector framework screen	57 cheap, profitable, under-followed names
Multi-agent scoring	5/28	5 specialists + adversarial reviewer + light model; calibrated on EXFY & SCOR	Worked, but judged to manufacture false precision → **retired**
Memo system & audit	5/29	Buffett-voice guide; 11 one-page memos; numbers-verification pass	Every memo had ≥1 wrong load-bearing number; lesson encoded
The Deep Flow	5/29–5/30	Full SOTP + DCF + Porter + sensitivity + VCP per name; agent-builds → lead-reviews	USNA, HCKT, AENT, MIND, SWAG, HOUR deep-flowed
ARTW build	5/30	Selected ARTW as an operator carve-out; deck + Excel model	Asset-protected, asymmetric search-acquisition thesis
Deck iteration	5/30–6/1	Voice, cold-reader audits, bottom-up value model, risks, AI disclosure	The current deck
Part 2	5/31	AI market-mapping: project plan + vendor stack	In progress

Session 1: Project Setup & Initial Screening (2026-05-26)

PROMPT 1: PROJECT INITIALIZATION & COMPANY SCREENING

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Set up project structure, identify qualifying companies ($10M-$500M EV on NYSE/Nasdaq), cross-reference with Aviv's sector expertise **Approach:** Parallel subagent architecture — one agent scanned Life_Admin project for background context, another researched micro/nano-cap universe **Key Outputs:**

- Identified ~1,500-2,500 companies in the $10M-$500M EV range across all sectors
- Narrowed to ~200-400 in technology/software/services
- Shortlisted 13 specific candidates across tiers (Tier 1: MNDO, DH, MCHX, DOMO, EGAN, SMSI, IDN, LPSN, VERI, GVP)
- Established router-based project documentation architecture **Where AI helped:** Rapidly synthesizing company screening across multiple sources, cross-referencing sector fit with professional background **Where AI may have gaps:** EV figures are approximate and may be stale; need real-time verification via SEC filings or financial data providers

Session 2: Data Pipeline & Framework (2026-05-27)

PROMPT 2: BUILD DATA PIPELINE

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Replace web-sourced data with reproducible Python pipeline pulling from yfinance and SEC EDGAR **Key Outputs:**

- `scripts/screen_companies.py` for yfinance-based screening (EV, ratios, FCF, analyst coverage)
- `scripts/fetch_edgar.py` for SEC filings + XBRL structured financials
- Pulled real-time data for the original 11 candidates; revealed EDGAR data current as of Q1 2026 **Where AI helped:** Fast scaffolding of reproducible data scripts with proper SEC User-Agent compliance **Where AI may have gaps:** Initial candidate list was based on stale web data — confirmed the need for framework-first approach

PROMPT 3: SELECTION FRAMEWORK DEVELOPMENT

Tool: Subagent (general-purpose research) **Purpose:** Build rigorous framework for interpreting Lunchline's criteria ("messy, mispriced, under-followed; avoid obvious names") **Key Outputs:**

- 10 messiness archetypes ranked by analytical richness
- 8 mispricing signal patterns specific to micro/nano-cap

- Quantitative under-followed thresholds
- 6-criterion weighted scoring framework (Value Creation 25%, Complexity 20%, Sector Fit 15%, Data 15%, Contrarian 15%, PE Realism 10%)
- Quantitative screening filters (primary + secondary)
- Research on what PE/search fund evaluators reward **Where AI helped:** Synthesized PE/search fund evaluation criteria into a defensible, scoreable framework **Where AI may have gaps:** Pre-scored old candidates using web data; those scores are now obsolete given fresh framework-driven screening

Session 3: Systematic Screening & Validation (2026-05-27 to 2026-05-28)

PROMPT 4: SYSTEMATIC FRAMEWORK SCREEN

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Apply framework filters to full NYSE/Nasdaq Tech + Comm Services universe **Approach:** Built `scripts/framework_screen.py` with three steps: Finviz universe pull, yfinance enrichment, framework filter application **Key Outputs:**

- 622 → 292 → 227 → 171 qualifying companies through funnel
- 83 in sweet spot (EV/Rev < 1.5x AND 0-3 analysts)
- 45 in ultra sweet spot ($30-250M EV + sweet spot)
- Industry distribution: 58 Software-Application, 34 Software-Infrastructure, 19 Internet Content, 13 IT Services, 10 Advertising, 9 Gaming **Where AI helped:** Built reproducible, multi-step screening script in one pass; handled Finviz pagination, yfinance batching, and filter application **Where AI may have gaps:** Some Chinese ADRs (CCG, LZMH, LGCL) show extreme EV/Rev ratios — may be data quality issues; need manual review

PROMPT 5: SANITY CHECK SWEET-SPOT CANDIDATES

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Validate framework by manually checking 2 candidates against all framework dimensions **Key Outputs:**

- GDEV (gaming): 6/7 checks passed. Profitable, cash-rich, under-followed.
- EXFY (SaaS): 7/7 checks passed. Cash-rich, FCF positive, classic broken-SaaS story. **Where AI helped:** Framework checks ran cleanly against real data; framework is producing legitimate candidates

Session 4: FinRobot Repo Review + Analyzer/Valuation Port (2026-05-28)

PROMPT 6: FINROBOT REPO SYNTHESIS

Tool: Subagent (general-purpose research) **Purpose:** Decide whether to install FinRobot (https://github.com/AI4Finance-Foundation/FinRobotSMARTPANTS_PRESERVED_283) or borrow specific methods. Initial framing rejected wholesale install (AutoGen overhead, OpenAI cost, originality risk for the case study). **Key Outputs:**

- Identified `finrobot_equity/` (Mar 2026) as the load-bearing module — 8 specialized agents mapping ~1:1 onto a buyside pitch.
- Confirmed README's "Adanos" reference is hype — no source code uses it. Real sentiment stack is Finnhub + Reddit + FinNLP.
- Surfaced three borrowable artifacts: (a) six tightly-scoped 10-K analyzer prompts pairing table + section + instruction, (b) FMP endpoints for analyst targets + competitor comps, (c) pure-Python DCF + football-field generator in `valuation_engine.py`.
- Flagged red flags: AutoGen group-chat orchestration unnecessary for one-shot analysis, 11-agent quant factor team is mostly stub system messages, notebook-driven dev in places. **Where AI helped:** Distinguished marketing claims from working code by reading the actual source tree; identified the small subset (~3 modules) worth porting vs. the bulk of the framework to skip. **Where AI may have gaps:** Did not verify FMP free-tier limits empirically; price-target endpoint is /v4 which typically requires the Starter plan.

PROMPT 7: PORT 10-K ANALYZER PROMPTS

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Port FinRobot's seven `analyzer.py` prompts (income / balance / cash flow / segment / business highlights / company description / risk) into `scripts/analyze_10k.py`. Each prompt pairs a financial table (yfinance) with the matching 10-K section (local EDGAR HTML) and an instruction. **Approach:** Wrote BS4 + regex section extractor with a TOC-discrimination heuristic — score each "Item N" candidate by how many other Item markers appear in the next 2000 chars (TOC entries cluster; real section starts are isolated). **Key Outputs:**

- `scripts/analyze_10k.py` — generates seven ready-to-paste prompt files per ticker.
- Verified against EXFY's 2026 10-K: Item 1 (54k chars), Item 1A (120k cap), Item 7 (48k). All three start at the actual section heading, not the TOC. **Where AI helped:** Wrote a pragmatic section extractor that handles the common iXBRL 10-K layout in one pass. **Where AI may have gaps:** Heuristic is tuned to recent 10-K HTML structure; older filings (pre-2020) may need different handling.

PROMPT 8: PORT FMP FETCHER

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Port FinRobot's `fmp_utils.py` endpoints into `scripts/fetch_fmp.py` — analyst price targets (consensus min/max/median), historical market cap, BVPS, single-company financial metrics, and competitor comps. **Approach:** Fixed a bug in the original `Revenue Growth` calc (uses `year_offset-1` when `year_offset==0`, which wraps to the oldest year — wrong). Added graceful error handling for missing API key, plan-tier limits, and bad symbols. Output paths under `data/fmp/<ticker>/`. **Where AI helped:** Identified and fixed the off-by-one growth-rate bug while porting. **Where AI may have gaps:** Did not exercise live API calls (requires FMP_API_KEY). The /v4 price-target endpoint likely requires the paid Starter plan; tested only the no-key error path locally.

PROMPT 10: PRE-EBITDA VALUATION METHODS RESEARCH

Tool: Subagent (general-purpose research) **Purpose:** Determine standard institutional practice for valuing companies with negative TTM EBITDA. Triggered by `valuation.py` aborting on EXFY (the framework's flagship 7/7 candidate) — needed to know what professional analysts actually do instead of EV/EBITDA + DCF. **Key Outputs:**

- Triangulation stack: EV/FCF (bypasses SBC debate) + Rule of 40-anchored EV/Revenue (Bessemer/Meritech) + Path-to-Profitability DCF with survival haircut (Damodaran). Cross-check with normalized year-3 EBITDA at search-fund 4-6x for downside anchor.
- The SBC fault line: value/PE camp treats SBC as cash cost; growth-equity adds it back. Lunchline's PE lens sides with value camp — should be stated explicitly in the deck.
- Meritech standard is EV/NTM Gross Profit, not EV/Revenue — rationale that GP captures business-model quality (vertical SaaS ~11x GP vs horizontal ~5x).
- Sourced from Damodaran (NYU Stern), Morningstar methodology, Bessemer Cloud Index, Meritech benchmarking, FTI Consulting / Oldfield (SBC critique). **Where AI helped:** Surfaced the SBC fault line as a methodological signal worth stating explicitly in the pitch, and the convergent "three-layer triangulation" framing used across Bessemer / Meritech / Damodaran. **Where AI may have gaps:** Multiple ranges cited are 2024-Q4 2025 data; sector defaults should be re-verified against current cloud-index data closer to pitch date.

PROMPT 12: PER-COMPANY PEER BENCHMARKS REFACTOR

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Replace hardcoded sector-generic multiples in `valuation.py` with per-company bespoke peer benchmarks persisted to CSV. Triggered by user feedback: "we shouldn't hard-code benchmarks without peer analysis" + "we should specify so that each company gets bespoke benchmarks" + "we can instruct the agents to update the csv whenever they do research so we don't have to re-do it every time." **Approach:** Designed `data/research/<ticker>/peer_benchmarks.csv` schema (14 columns: peer ticker/name/inclusion tag + 4 multiples + 4 margins/growth + date/source/notes). Each ticker gets its own file. Subject anchor row pre-populated from yfinance. When file missing, script writes the empty template with instructions and exits — code-level enforcement of methodology #3 ("comp-proven, not theoretical"). Added agent instruction to CLAUDE.md so future sessions know to persist peer findings. **Key Outputs:**

- Removed hardcoded R40 band table (1/2/3.5/5.5/7.5/10x). Replaced with `value_ev_revenue` (peer-derived multiple, R40 reported as diagnostic context).
- Three multiple-resolution sources: peer CSV (target=median, low=min, high=max across primaries) > CLI override > skip. Each method labels its source in output.
- Validation: EXFY with 3 primary peers → blended $5.74 (+397% vs $1.16). AAPL regression: unchanged when CLI overrides match prior defaults.
- New `TODO.md` capturing pending build/research items per the new methodology direction. **Where AI helped:** Designed a clean three-tier multiple resolution (peer → CLI → skip) that prevents silent defaults while keeping CLI bypass available for one-off runs. Schema includes the subject row as anchor so agents have local reference. **Where AI may have gaps:** Schema may need additional columns over time (e.g., size adjustment factor, growth-adjusted multiple, NTM vs LTM toggle). Defer until real peer-research workflows surface what's missing.

PROMPT 11: EXTEND VALUATION.PY WITH PRE-EBITDA METHODS

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Implement the four pre-EBITDA methods from Prompt 10's research into `scripts/valuation.py`. **Key Outputs:**

- Added `value_ev_fcf`, `value_ev_gp`, `value_rule40`, `value_path_dcf`.
- Method routing via `--methods auto` (positive-EBITDA → ev_ebitda+dcf, pre-EBITDA → ev_fcf+ev_gp+rule40+path_dcf) or explicit `--methods` list.
- Sensitivity matrix now auto-picks EBITDA or FCF mode.
- EXFY validation: 4 methods bracket $2.50-$4.92 → blended $3.86 vs $1.15 current (+236%). Consistent with net-cash position making effective EV/FCF ~2.4x trailing.
- AAPL regression: unchanged outputs from prior version ($173 EV/EBITDA, $110 DCF, $145 blended). **Where AI helped:** Carefully refactored without breaking the positive-EBITDA path (verified by AAPL regression). Auto-routing kept the CLI ergonomic despite adding 12 new flags. **Where AI may have gaps:** Path-DCF defaults (15/12/10/8/6% growth, -5/0/5/10/15% margin) are sector-generic; per-company tuning required before pitch defense.

PROMPT 9: PORT DCF + FOOTBALL FIELD

Tool: Claude Code (Claude Opus 4.7) **Purpose:** Port `valuation_engine.py` (10-yr DCF, EV/EBITDA, peer comp) and `sensitivity_analyzer.py` (2D revenue × margin matrix) into `scripts/valuation.py`. Pure numpy/pandas/matplotlib — no LLM dependency. **Approach:** Exposed all DCF assumptions as CLI flags (`--growth-near`, `--growth-far`, `--terminal-growth`, `--wacc`, `--fcf-conversion`) rather than hardcoding FinRobot's defaults, so assumptions are visible and iterable. Bear/bull range = ±200bps WACC, ±50bps terminal growth. Confidence-weighted blended target reported. **Key Outputs:**

- Sanity-tested on AAPL: EV/EBITDA @ 18x → $173 (low $134, high $213), DCF @ 5/3% growth + 8.5% WACC → $110 (low $78, high $186). Blended $145 vs. current $311.46 — i.e. the model says AAPL is expensive on these assumptions, which is the expected directional result.
- Outputs: `dcf_projection.csv`, `sensitivity.csv`, `football_field.png`, `summary.json`.
- Pre-EBITDA companies (e.g. EXFY current TTM) trigger an early-exit warning recommending EV/Sales or SOTP — DCF is uninformative when base FCF is negative. **Where AI helped:** Surfaced that the framework's many sweet-spot candidates are pre-EBITDA, which forces a methodological choice (multi-method valuation per-candidate vs. one method for all). **Where AI may have gaps:** Football-field chart is matplotlib-default styling — fine for working artifacts, may want a designer pass before the final deck.

Session 6: Memo system, candidate expansion, and a numbers audit (2026-05-29)

Tool: Claude Code (Claude Opus 4.8, 1M context)

PROMPT: BUILD AN INVESTMENT-MEMO VOICE GUIDE FROM EXEMPLARS

Purpose: The project had a design system but no *writing* guidance; Aviv cited Buffett's letters as the reference. Dispatched a research subagent over Berkshire letters, Howard Marks memos, and an Einhorn writeup. **Output:** `docs/memo-voice.md` (8 voice principles with verbatim exemplar sentences, plain-English vocabulary table, inline-sourcing patterns, narrative-arc templates, annotated 1-page skeleton) + `docs/memo-principles.md` (checklist + forbidden moves). **Where AI helped:** Extracted concrete, imitable patterns rather than generic "write clearly" advice. **Where AI may have gaps:** Two source PDFs were binary-locked and substituted; noted in the doc.

PROMPT: TRIAGE THE 21 UNEVALUATED LAYER-2 CANDIDATES, THEN BUILD PITCH MEMOS

Approach: A batch script ran dossier + voting + deal-status on all 21 (16 survived); 8 parallel subagents produced findings on the 16; 10 parallel subagents drafted one-page HTML memos against the template. 11 memos total in `mockups/pitches/`. **Where AI helped:** Parallel fan-out turned a multi-day sweep into a single session. **Where AI HURT — important disclosure:** the memo-drafting agents inherited unverified numbers from the findings stage and, in places, fabricated plausible figures. A follow-up audit (below) found at least one wrong load-bearing number in **every** memo.

PROMPT: VERIFY EVERY MEMO'S NUMBERS AGAINST FILINGS (THE AUDIT)

Purpose: After catching a 7x adjusted-EBITDA error in IZEA by hand, ran one verification subagent per memo — GAAP vs XBRL, non-GAAP vs the 10-K reconciliation table, derived figures recomputed. **Key outputs:** Four repeatable failure modes documented (data-vendor aggregate fields, non-GAAP inherited not sourced, conflation of similar-magnitude figures, period/share-count slips). Thesis-altering corrections on CDLX, MIND, SCOR, IZEA, UONE; RSSS/UPLD/SWAG/AENT survived clean. Lesson encoded in `docs/memo-principles.md`. **Where AI helped:** The verification pass is the most valuable single step in the session — it caught errors that read as plausible and would have survived a human read-through. **Where AI may have gaps / lesson:** LLM agents will state a confident specific number that is wrong; treat any agent-produced figure as unverified until checked against the primary filing. yfinance `enterpriseValue` / `marketCap` / `freeCashFlow` are especially unreliable for dual-class/earnout/preferred structures — compute from the balance sheet instead.

Sessions 7–13 — summarized from transcripts (2026-05-28 → 2026-06-01)

The entries below are condensed (representative prompts + outcomes), not per-prompt verbatim.

Session 7: Multi-agent scoring calibration (2026-05-28)

Tool: Claude Code (Opus) + parallel subagents **Representative prompts:** "Compare a 7/7 with a 3/7"; "how would we run the weighted criterion scoring?"; "split into multiple subagents for greater task-focus and objectivity?"; "ignore my own expertise at this phase… add an adversarial 7th pass… using solid investment logic"; then the per-dimension specialist prompts and the adversarial-reviewer prompts for **EXFY** and **SCOR**. **What ran:** A scoring pipeline — 5 dimension specialists (messiness, value-creation, data, contrarianism, PE-realism) + an adversarial reviewer + a lightweight returns model — plus voting-structure and deal-status detectors. Calibrated on EXFY (rescaled 3.35/5; detector caught an 84.7% voting-trust deal-breaker) and SCOR (the adversarial pass caught that a reported $103.9M "net income" was a non-cash recap gain, plus a covenant waiver — devastating). **Where AI helped:** The adversarial reviewer and the voting/deal-status detectors caught structural deal-breakers (super-voting trusts, mid-deal targets like LPSN/KPLT) the specialists missed. **Where AI fell short:** The numeric scoring produced **false precision** — confident 1–5 scores that didn't survive scrutiny. This directly motivated retiring the scoring approach (Session 10).

Session 8: Pipeline consolidation & cleanup (2026-05-28)

Representative prompts: "review what we've done… any files we should clean up that are now redundant? duplications? is claude.md clean and routing clearly?"; "double-check each part is safe to delete… don't delete reports"; "the pipeline shouldn't be in claude.md, route to a separate file"; "commit+push". **What ran:** De-duplicated docs/workflows, moved the pipeline spec out of CLAUDE.md into its own routed doc, committed. Ran findings/specialist passes on CDLX, NRDY, FAST (diagnostic), THRY. **Where AI helped:** Fast, careful repo hygiene with safety checks before deletes.

Session 9: Memo voice, one-page memos & the numbers audit (2026-05-29)

(Detailed in Session 6 above — Buffett/Marks/Einhorn voice guide, 11 one-page HTML memos, then a per-memo verification pass that found a wrong load-bearing number in every memo.) The key disclosure: **memo-drafting agents fabricated plausible figures**; only the audit caught them.

Session 10: Retiring scoring → the Deep Flow & candidate tournament (2026-05-29 → 05-30)

Representative prompts: "How are we picking from the 42?"; "is there an overall SaaS valuation decline due to AI?" (RSSS); the skeptical-analyst **USNA** deep-flow prompt, then "I'm the reviewing lead; we iterate until we've extracted everything the filings allow"; "continue analyzing more companies using the deep flow"; the **AENT** deep-flow prompt; "why did the stock collapse 54%… did we look at the earnings calls?"; "this model's EV is the stock is worth $12.76 vs $11.61?"; "help me pick a new stock we haven't touched out of the 57"; "enrich the catalog with owner-dependence / key-man / acquirability signals" (with the correction: "don't override our previous methodology — enrich it"). **What ran:** The 6-criterion scoring was **retired** (it had called USNA a value trap; the Deep Flow reversed then disciplined that). Full Deep Flows (SOTP + driver model + DCF on cash taxes + Porter + sensitivity + VCP + kill criteria) on USNA, HCKT, AENT, MIND, SWAG, HOUR. Catalog enriched with search-fund acquirability signals. Industry-structure research on coffee (JVA), footwear (WEYS), ACU. **Where AI helped:** The agent-builds → lead-reviews loop extracted far more from the filings than a single pass; the USNA reversal showed the depth was worth it. **Where AI fell short:** Needed the lead to repeatedly re-anchor it (e.g., "don't override the methodology, enrich it") — without a method in mind, the volume of analysis drifts.

Session 11: ARTW selection, deep flow & first build (2026-05-30)

Representative prompts: "Tell me about ARTW"; "let's do our deep dive analysis on this"; "explain the demand/conversion math — do we have actual metrics on whether there was demand they couldn't meet?"; the filings-comb prompt for demand-constrained evidence in the modular segment; "let's build the slide deck and the excel model"; then the workbook-builder and deck-builder agent prompts. **What ran:** ARTW chosen as a take-private-and-operate carve-out. Combed filings for the demand-constraint signal (backlog +103%, thin commercial org). Built `ARTW model.xlsx` and the Part-1 HTML deck. **Where AI helped:** Surfaced ARTW from the untouched screen and assembled the carve-out package quickly.

Session 12: Deck voice iteration + portability (2026-05-30 → 05-31)

Representative prompts: Feedback "did you use the Buffett style? it feels too salesy… more concise, more to-the-point, avoid marketing language"; the no-context cold-reader audit prompts; "rewrite for voice and clarity, keep the design system"; "no need to explain concepts in-line like 'operating margin' — assume savvy readers"; "you reintroduced em-dashes — remove them"; slide-specific line edits; gather company facts each with an active source URL; then "log to changelog… continue from a new session"; git-ignore questions to get ARTW data + the screening catalog onto the laptop; "everything on remote?". **What ran:** Multiple voice rewrites verified by cold-reader subagents; em-dash purges; sourcing pass; un-ignored the research/catalog for portability; commit+push. **Where AI helped:** Cold-reader audit subagents gave honest outside-eye feedback on clarity. **Where AI fell short:** Repeatedly reintroduced banned patterns (em-dashes, salesy phrasing, finance-101 glosses) until given explicit examples of "good."

Session 13: Bottom-up value model + deck polish + this log (2026-05-31 → 06-01)

Representative prompts: rebase the value-creation waterfall to the control basis and the model; re-judge the monitoring lever ("I don't agree with your take… the question is upsell + industry structure, not whether someone's already doing it"); "model the operational path bottom-up… these compound with the multiple… don't cap it top-down"; "critique the approach and help me decide if we should use it at all"; tighten the risks slide ("reframe 'what limits it' as mitigations… really concise"); make the AI disclosure concise (three sections); "use Motif to get the real transcripts and produce a summarized prompt report." **What ran:** Rebased the waterfall ($14.0M → $22.9M, control basis); an independent agent reversed the dismissive monitoring take (→ installed-base annuity, ~$7B fragmented services market); built a **benchmarked bottom-up value-creation model** (sales engine, up-market mix, monitoring, multiple compounding) in `model.md §v4`, then — after a self-critique flagged it as an assumption-stacked hero-bridge — **disciplined it back** for the deck (base-anchored, re-rate de-emphasized). Rebuilt the risks slide around the value-creation levers; rewrote this AI-disclosure appendix; swept em-dashes; pulled the transcripts via Motif to reconstruct this log. **Where AI helped:** Independent research agents (monitoring industry structure, sales-effectiveness benchmarks) produced sourced driver assumptions; a self-critique caught the bottom-up model overreaching before it shipped. **Where AI fell short:** The first monitoring pass over-indexed on "a competitor already does it" and lost the bigger upsell/industry-structure question — corrected only after the lead pushed back. Pattern confirmed: AI needs strict control, iteration, and examples of "what good looks like" to be net-positive.